

Analysis of Heart Disease

June 2, 2020

0.1 Introduction

Diagnosing heart disease is an important part of a doctor's job. This allows them to take necessary steps in order to prevent severe health degradation in the patient. This is not always possible because it is hard to have a physician to diagnose the person with heart disease all the time. This notebook is an analysis of a few classification algorithms that can be used to diagnose heart disease automatically. Classification algorithms can possibly be implemented in smart medical devices that can alert the patient of possible heart disease which allows them to seek help to mitigate the health risks.

```
[1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import GridSearchCV, train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.svm import SVC
from sklearn.ensemble import RandomForestClassifier, AdaBoostClassifier
from sklearn.metrics import confusion_matrix, classification_report, roc_curve, \
    auc
import pickle
import warnings
import matplotlib.style as style
style.use("seaborn-whitegrid")
warnings.filterwarnings('ignore')
```

0.2 Exploratory Data Analysis

From the description of the data we know

age - age in years

sex - (1 = male; 0 = female)

cp - chest pain type

trestbps - resting blood pressure (in mm Hg on admission to the hospital)

chol - serum cholestoral in mg/dl

fbs - (fasting blood sugar > 120 mg/dl) (1 = true; 0 = false)

restecg - resting electrocardiographic results
 thalach - maximum heart rate achieved
 exang - exercise induced angina (1 = yes; 0 = no)
 oldpeak - ST depression induced by exercise relative to rest
 slope - the slope of the peak exercise ST segment
 ca - number of major vessels (0-3) colored by flourosopy
 thal - 3 = normal; 6 = fixed defect; 7 = reversable defect
 target - have disease or not (1=yes, 0=no)

```
[2]: # Load the data
data = pd.read_csv("heart.csv")
data.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 303 entries, 0 to 302
Data columns (total 14 columns):
#   Column      Non-Null Count  Dtype
---  -
0   age         303 non-null    int64
1   sex         303 non-null    int64
2   cp          303 non-null    int64
3   trestbps    303 non-null    int64
4   chol        303 non-null    int64
5   fbs         303 non-null    int64
6   restecg     303 non-null    int64
7   thalach     303 non-null    int64
8   exang       303 non-null    int64
9   oldpeak     303 non-null    float64
10  slope       303 non-null    int64
11  ca          303 non-null    int64
12  thal        303 non-null    int64
13  target      303 non-null    int64
dtypes: float64(1), int64(13)
memory usage: 33.3 KB
```

From the loaded data we can see that there are 303 observations with 14 features. There are no null values as shown in the non-null count. This means we don't need to deal with non-null values in this dataset.

```
[3]: data.head()

[3]:   age  sex  cp  trestbps  chol  fbs  restecg  thalach  exang  oldpeak  slope  \
0   63   1   3     145    233   1         0     150     0       2.3     0
1   37   1   2     130    250   0         1     187     0       3.5     0
2   41   0   1     130    204   0         0     172     0       1.4     2
```

3	56	1	1	120	236	0	1	178	0	0.8	2
4	57	0	0	120	354	0	1	163	1	0.6	2

	ca	thal	target
0	0	1	1
1	0	2	1
2	0	2	1
3	0	2	1
4	0	2	1

```
[4]: data = data.drop_duplicates()
data.shape
```

```
[4]: (302, 14)
```

```
[5]: # Get summary statistics
data.describe()
```

```
[5]:
```

	age	sex	cp	trestbps	chol	fbs	\
count	302.000000	302.000000	302.000000	302.000000	302.000000	302.000000	
mean	54.42053	0.682119	0.963576	131.602649	246.500000	0.149007	
std	9.04797	0.466426	1.032044	17.563394	51.753489	0.356686	
min	29.00000	0.000000	0.000000	94.000000	126.000000	0.000000	
25%	48.00000	0.000000	0.000000	120.000000	211.000000	0.000000	
50%	55.50000	1.000000	1.000000	130.000000	240.500000	0.000000	
75%	61.00000	1.000000	2.000000	140.000000	274.750000	0.000000	
max	77.00000	1.000000	3.000000	200.000000	564.000000	1.000000	

	restecg	thalach	exang	oldpeak	slope	ca	\
count	302.000000	302.000000	302.000000	302.000000	302.000000	302.000000	
mean	0.526490	149.569536	0.327815	1.043046	1.397351	0.718543	
std	0.526027	22.903527	0.470196	1.161452	0.616274	1.006748	
min	0.000000	71.000000	0.000000	0.000000	0.000000	0.000000	
25%	0.000000	133.250000	0.000000	0.000000	1.000000	0.000000	
50%	1.000000	152.500000	0.000000	0.800000	1.000000	0.000000	
75%	1.000000	166.000000	1.000000	1.600000	2.000000	1.000000	
max	2.000000	202.000000	1.000000	6.200000	2.000000	4.000000	

	thal	target
count	302.000000	302.000000
mean	2.314570	0.543046
std	0.613026	0.498970
min	0.000000	0.000000
25%	2.000000	0.000000
50%	2.000000	1.000000
75%	3.000000	1.000000
max	3.000000	1.000000

```
[6]: target = data.target.apply(lambda x: "Healthy " if x==0 else "Heart Disease")

print("Total count of each type of patient")
print(data.target.value_counts())
```

Total count of each type of patient


1 164

0 138

Name: target, dtype: int64

```
[7]: # Plot the Frequency
sns.countplot(target);

plt.xlabel('Patient Status')
plt.ylabel('Frequency')
plt.title('Frequency of Diseased and Healthy Patients');
```

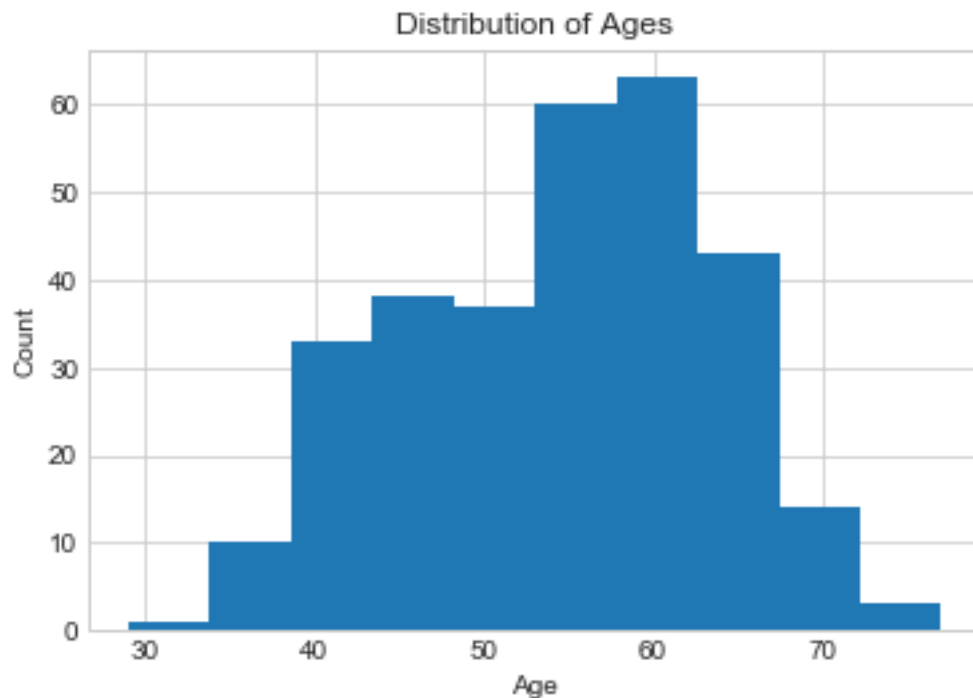


```
[8]: no_disease = len(data[data.target == 0])
disease = len(data[data.target == 1])
total = len(data)

print("Percentage of healthy patients: {:.2f}%".format(no_disease*100/total))
print("Percentage of patients with heart disease: {:.2f}%".format(disease*100/
    ↳total))
```

Percentage of healthy patients: 45.70%
Percentage of patients with heart disease: 54.30%

```
[9]: # See distribution of each column
plt.hist(x=data.age)
plt.title("Distribution of Ages")
plt.xlabel("Age")
plt.ylabel("Count");
```



From the histogram, we can see that most of the patients in the dataset are between the ages of 55 and 60. The histogram also shows that the patient's age is approximately normally distributed.

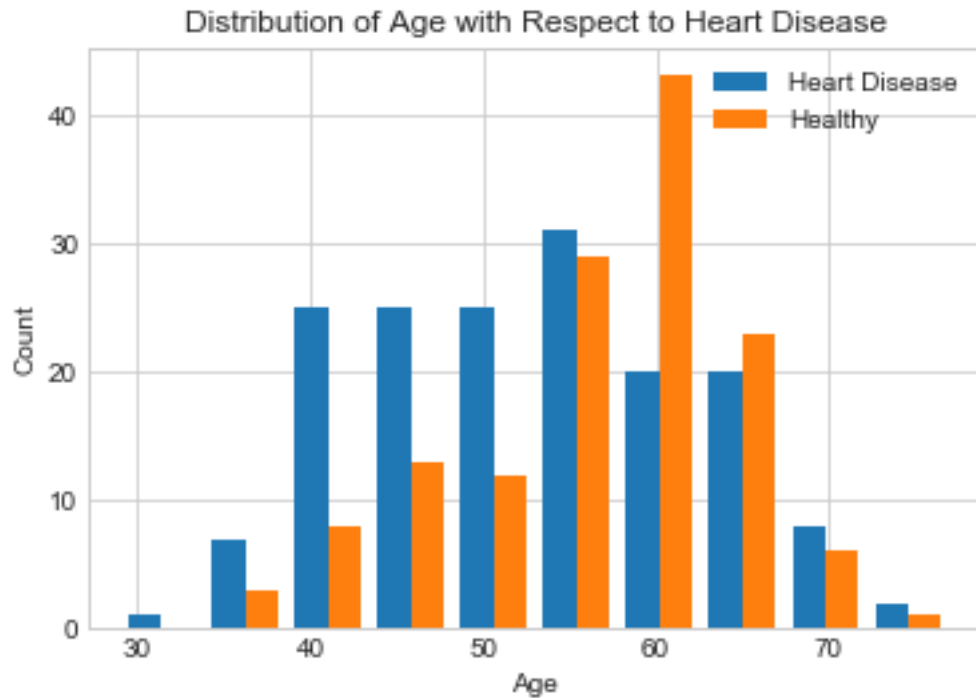
```
[10]: # Seperate the heart diseased age column
x1 = data.age[data.target == 1]

# Seperate the healthy age column
x2 = data.age[data.target == 0]

names = ["Heart Disease", "Healthy"]
plt.hist([x1, x2], bins = int(303/30), label=names)

# Plot formatting
plt.legend()
plt.xlabel('Age')
plt.ylabel('Count')
```

```
plt.title('Distribution of Age with Respect to Heart Disease');
```



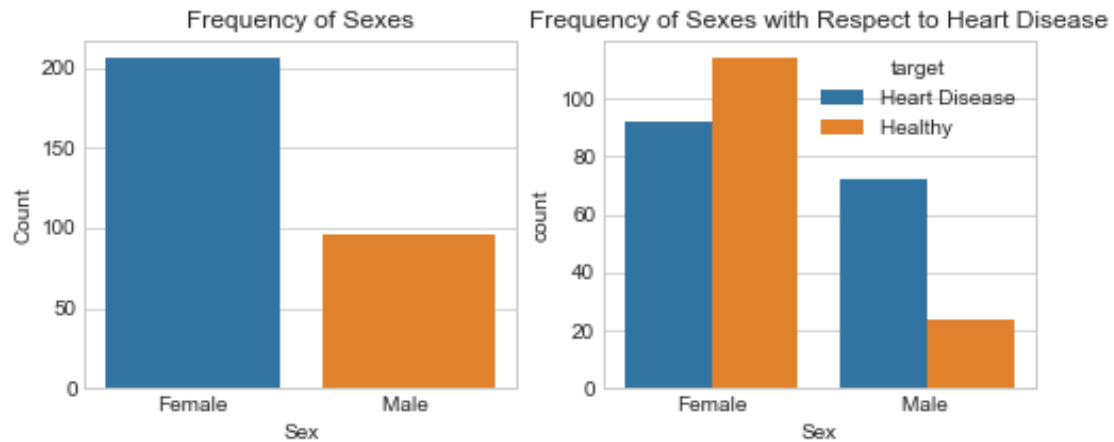
From the distribution we can see that heart disease seems to affect the age group of 50 to 60 the most. There is a sharp

```
[11]: sexes = data.sex.apply(lambda x: "Male " if x==0 else "Female")
target = data.target.apply(lambda x: "Healthy " if x==0 else "Heart Disease")

# Explore the distribution of sexes
fig, axs = plt.subplots(ncols=2)
fig.set_figheight(3)
fig.set_figwidth(8)

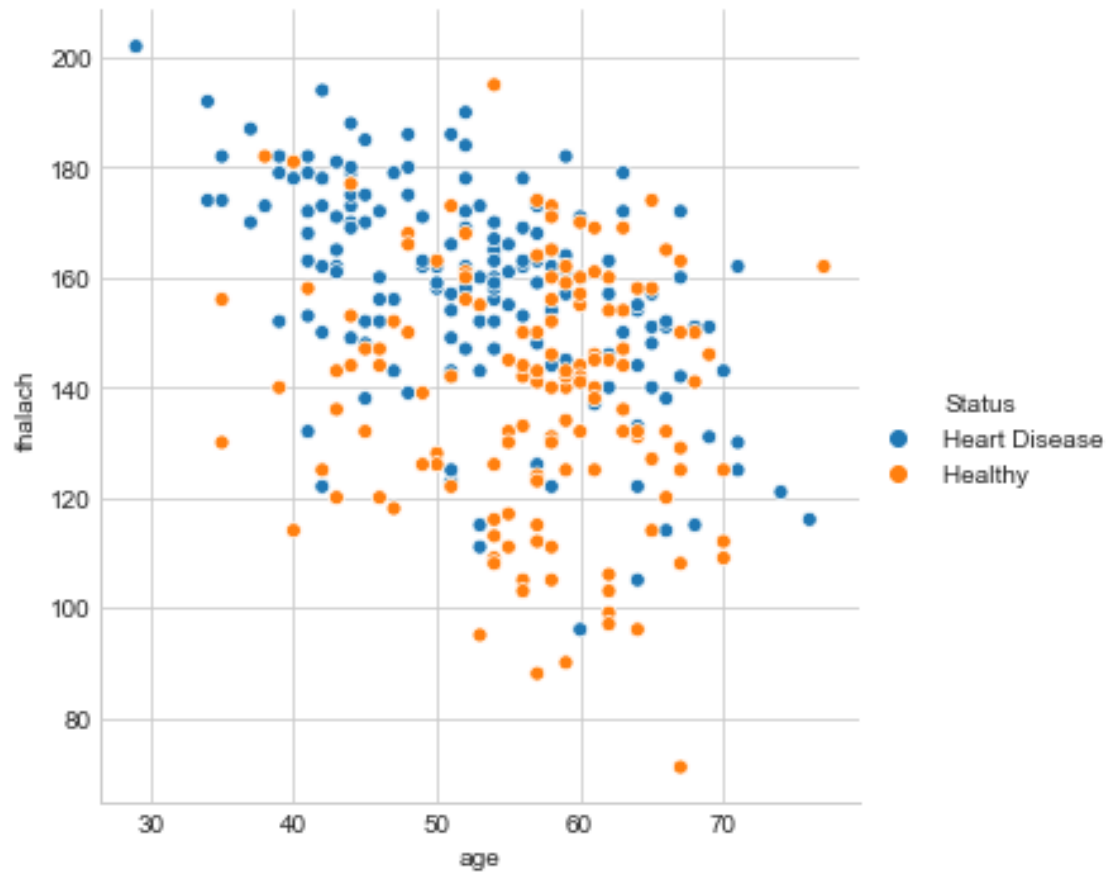
# Pl
sns.countplot(sexes, ax=axs[0])
axs[0].set_xlabel("Sex")
axs[0].set_ylabel("Count")
axs[0].set_title("Frequency of Sexes");

sns.countplot(sexes, hue=target, ax=axs[1])
axs[1].set_xlabel("Sex")
axs[1].set_title("Frequency of Sexes with Respect to Heart Disease");
```




More female patients were collected and there are more healthy females compared to females with heart disease. Males in the data are primarily have heart disease. There is a possibility the model might be biased towards to males being more likely to have heart disease.

```
[12]: Status = data.target.apply(lambda x: "Healthy " if x==0 else "Heart Disease")
      Status=Status.rename("Status")
      sns.relplot(x="age", y="thalach",hue=Status,data=data,legend="full");
```



Another interesting observation is that maximum heart rate decreases as the patient is older. Although maximum heart rate is on average higher for patients with heart disease compared to patients without heart disease.


```
[13]: plt.hist(x=data.oldpeak)
plt.title("Distribution of Oldpeak")
plt.xlabel("Oldpeak")
plt.ylabel("Count");
```

Seems like Oldpeak is not distributed normally which might cause problems with regression classifiers. Maybe the use of Box-Cox transformation can help.

```
[14]: target = data.target.apply(lambda x: "Healthy" if x==0 else "Heart Disease")
      exang = data.exang.apply(lambda x: "Induced Angina" if x==1 else "No Effect")

      sns.countplot(exang, hue=target)
      plt.title("Distribution of Exang")
      plt.xlabel("Exang")
      plt.ylabel("Count")
      plt.legend(title="Status");
```



Patients that were subject to an induced angina from the stress exercise tended to be healthy and less likely to have heart disease.

0.3 Feature Engineering

Indicator Values

```
[15]: for col in data.columns:
      if(col != "target"):
          print("-----{}-----".format(col))
          print(data[col].unique())
```

```
-----age-----
[63 37 41 56 57 44 52 54 48 49 64 58 50 66 43 69 59 42 61 40 71 51 65 53
 46 45 39 47 62 34 35 29 55 60 67 68 74 76 70 38 77]
-----sex-----
[1 0]
-----cp-----
[3 2 1 0]
-----trestbps-----
[145 130 120 140 172 150 110 135 160 105 125 142 155 104 138 128 108 134
 122 115 118 100 124 94 112 102 152 101 132 148 178 129 180 136 126 106
 156 170 146 117 200 165 174 192 144 123 154 114 164]
-----chol-----
[233 250 204 236 354 192 294 263 199 168 239 275 266 211 283 219 340 226
 247 234 243 302 212 175 417 197 198 177 273 213 304 232 269 360 308 245]
```

```

208 264 321 325 235 257 216 256 231 141 252 201 222 260 182 303 265 309
186 203 183 220 209 258 227 261 221 205 240 318 298 564 277 214 248 255
207 223 288 160 394 315 246 244 270 195 196 254 126 313 262 215 193 271
268 267 210 295 306 178 242 180 228 149 278 253 342 157 286 229 284 224
206 167 230 335 276 353 225 330 290 172 305 188 282 185 326 274 164 307
249 341 407 217 174 281 289 322 299 300 293 184 409 259 200 327 237 218
319 166 311 169 187 176 241 131]
-----fbs-----
[1 0]
-----restecg-----
[0 1 2]
-----thalach-----
[150 187 172 178 163 148 153 173 162 174 160 139 171 144 158 114 151 161
 179 137 157 123 152 168 140 188 125 170 165 142 180 143 182 156 115 149
 146 175 186 185 159 130 190 132 147 154 202 166 164 184 122 169 138 111
 145 194 131 133 155 167 192 121 96 126 105 181 116 108 129 120 112 128
 109 113 99 177 141 136 97 127 103 124 88 195 106 95 117 71 118 134
 90]
-----exang-----
[0 1]
-----oldpeak-----
[2.3 3.5 1.4 0.8 0.6 0.4 1.3 0. 0.5 1.6 1.2 0.2 1.8 1. 2.6 1.5 3. 2.4
 0.1 1.9 4.2 1.1 2. 0.7 0.3 0.9 3.6 3.1 3.2 2.5 2.2 2.8 3.4 6.2 4. 5.6
 2.9 2.1 3.8 4.4]
-----slope-----
[0 2 1]
-----ca-----
[0 2 1 3 4]
-----thal-----
[1 2 3 0]

```

From reading the description and investigating the unique values, we can conclude that thal, ca, slope, exang, restecg, fbs, cp and sex are indicator features. The features thal, slope, restecg, and cp need to be transformed into dummy variables.

```

[16]: # Get the dummies
cp = pd.get_dummies(data.cp, prefix="cp")
restecg = pd.get_dummies(data.restecg, prefix="restecg")
slope = pd.get_dummies(data.slope, prefix="slope")
ca = pd.get_dummies(data.ca, prefix="ca")
thal = pd.get_dummies(data.thal, prefix="thal")

# Get all columns except the last one
cp = cp.iloc[:, 1:cp.shape[1]]
restecg = restecg.iloc[:, 1:restecg.shape[1]]
slope = slope.iloc[:, 1:slope.shape[1]]
ca = ca.iloc[:, 1:ca.shape[1]]
thal = thal.iloc[:, 1:thal.shape[1]]

```

```
# Verify that the columns were made into dummy variables
for dummy in cp,restecg,slope,ca,thal:
    print(dummy.columns)
```

```
Index(['cp_1', 'cp_2', 'cp_3'], dtype='object')
Index(['restecg_1', 'restecg_2'], dtype='object')
Index(['slope_1', 'slope_2'], dtype='object')
Index(['ca_1', 'ca_2', 'ca_3', 'ca_4'], dtype='object')
Index(['thal_1', 'thal_2', 'thal_3'], dtype='object')
```

```
[17]: # Seperate target value and features
X = data.loc[:, data.columns != 'target']
y = data.target

print("Shape before removal {}".format(X.shape))

# Remove the categorical columns
for cat in ['cp', 'restecg', 'slope', 'ca', 'thal']:
    X = X.loc[:, X.columns != cat]

# Verify 5 columns were removed
print("Shape after removal {}".format(X.shape))
```

```
Shape before removal (302, 13)
Shape after removal (302, 8)
```

```
[18]: # Add the dummy variable columns
X = pd.concat([X, cp,restecg,slope,ca,thal], axis=1)
X.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 302 entries, 0 to 302
Data columns (total 22 columns):
#   Column      Non-Null Count  Dtype
---  -
0   age         302 non-null   int64
1   sex         302 non-null   int64
2   trestbps    302 non-null   int64
3   chol        302 non-null   int64
4   fbs         302 non-null   int64
5   thalach     302 non-null   int64
6   exang       302 non-null   int64
7   oldpeak     302 non-null   float64
8   cp_1        302 non-null   uint8
9   cp_2        302 non-null   uint8
10  cp_3        302 non-null   uint8
11  restecg_1   302 non-null   uint8
```

```

12 restecg_2  302 non-null  uint8
13 slope_1    302 non-null  uint8
14 slope_2    302 non-null  uint8
15 ca_1       302 non-null  uint8
16 ca_2       302 non-null  uint8
17 ca_3       302 non-null  uint8
18 ca_4       302 non-null  uint8
19 thal_1     302 non-null  uint8
20 thal_2     302 non-null  uint8
21 thal_3     302 non-null  uint8
dtypes: float64(1), int64(7), uint8(14)
memory usage: 35.4 KB

```

```

[19]: # Scale the non-categorical variables
scaled_features = data.copy()
col_names = ['oldpeak', 'thalach', 'chol', 'trestbps', 'age']
features = scaled_features[col_names]
scaler = StandardScaler().fit(features.values)
features = scaler.transform(features.values)

X[col_names] = features
X.describe()

```

```

[19]:
count    age          sex      trestbps      chol      fbs  \
mean    -2.724090e-16   0.682119 -8.053712e-16 -2.086263e-17  0.149007
std      1.001660e+00   0.466426  1.001660e+00  1.001660e+00  0.356686
min     -2.814192e+00   0.000000 -2.144521e+00 -2.332210e+00  0.000000
25%     -7.107878e-01   0.000000 -6.617119e-01 -6.870826e-01  0.000000
50%      1.195033e-01   1.000000 -9.140084e-02 -1.161266e-01  0.000000
75%      7.283833e-01   1.000000  4.789102e-01  5.467629e-01  0.000000
max      2.499671e+00   1.000000  3.900776e+00  6.145034e+00  1.000000

count    thalach    exang    oldpeak    cp_1    cp_2  ...  \
mean    -4.087974e-16  0.327815 -1.948405e-16  0.165563  0.284768  ...
std      1.001660e+00  0.470196  1.001660e+00  0.372305  0.452053  ...
min     -3.436149e+00  0.000000 -8.995441e-01  0.000000  0.000000  ...
25%     -7.137164e-01  0.000000 -8.995441e-01  0.000000  0.000000  ...
50%      1.281605e-01  0.000000 -2.096081e-01  0.000000  0.000000  ...
75%      7.185677e-01  1.000000  4.803280e-01  0.000000  1.000000  ...
max      2.292987e+00  1.000000  4.447460e+00  1.000000  1.000000  ...

count    restecg_2    slope_1    slope_2    ca_1    ca_2    ca_3  \
mean      0.013245    0.463576    0.466887    0.215232    0.125828    0.066225
std      0.114512    0.499499    0.499730    0.411665    0.332206    0.249088

```

min	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000
25%	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000
50%	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000
75%	0.000000	1.000000	1.000000	0.000000	0.000000	0.000000
max	1.000000	1.000000	1.000000	1.000000	1.000000	1.000000

	ca_4	thal_1	thal_2	thal_3
count	302.000000	302.000000	302.000000	302.000000
mean	0.013245	0.059603	0.546358	0.387417
std	0.114512	0.237142	0.498673	0.487969
min	0.000000	0.000000	0.000000	0.000000
25%	0.000000	0.000000	0.000000	0.000000
50%	0.000000	0.000000	1.000000	0.000000
75%	0.000000	0.000000	1.000000	1.000000
max	1.000000	1.000000	1.000000	1.000000

[8 rows x 22 columns]

0.4 Model Building

```
[20]: # 80-20 split of train test data
X_train, X_test, y_train, y_test = train_test_split(X,y, test_size=0.2,
↳random_state=43)

print("The size of the training set is {} rows".format(X_train.shape[0]))
print("The size of the test set is {} rows".format(X_test.shape[0]))
```

The size of the training set is 241 rows

The size of the test set is 61 rows

```
[21]: param_grid = {
    'penalty':['l1', 'l2', 'elasticnet', 'none'],
    'l1_ratio':np.linspace(0,1, num=10),
    'random_state': [123]
}
# instantiate the grid
grid = GridSearchCV(LogisticRegression(max_iter=1000), param_grid, cv=10,
↳scoring='accuracy', return_train_score=True)
logreg = grid.fit(X_train,y_train)
```

```
[22]: logreg.best_estimator_
```

```
[22]: LogisticRegression(C=1.0, class_weight=None, dual=False, fit_intercept=True,
    intercept_scaling=1, l1_ratio=0.0, max_iter=1000,
    multi_class='auto', n_jobs=None, penalty='l2',
    random_state=123, solver='lbfgs', tol=0.0001, verbose=0,
    warm_start=False)
```

```
[23]: def plot_confusion_matrix(X_test,y_test, estimator,name="Estimator"):
    cm = confusion_matrix(y_test, estimator.predict(X_test))
    sns.heatmap(cm,annot=True)
    plt.ylabel('True label')
    plt.xlabel('Predicted label')
    plt.title("Confusion matrix of {}".format(name))
    plt.show()

def plot_roc(X_test,y_test,estimator,name="Estimator"):
    probs = estimator.predict_proba(X_test)
    preds = probs[:,1]
    fpr, tpr, threshold = roc_curve(y_test, preds)
    roc_auc = auc(fpr, tpr)

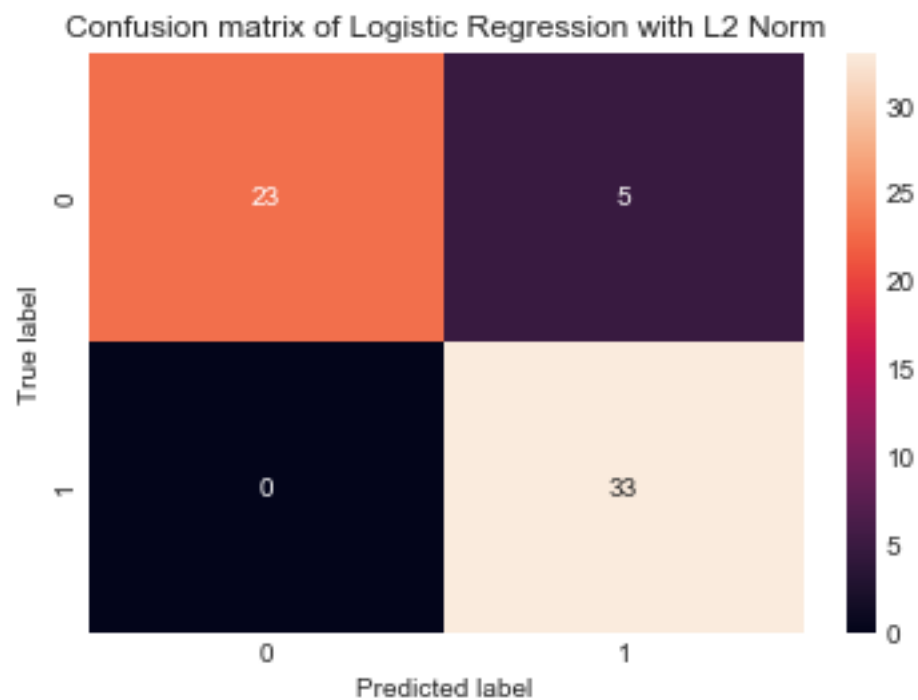
    plt.title('Receiver Operating Characteristic of {}'.format(name))
    plt.plot(fpr, tpr, 'b', label = 'AUC = %0.2f' % roc_auc)
    plt.legend(loc = 'lower right')
    plt.plot([0, 1], [0, 1], 'r--')
    plt.xlim([0, 1])
    plt.ylim([0, 1])
    plt.ylabel('True Positive Rate')
    plt.xlabel('False Positive Rate')
    plt.show()

def display_metrics(X_test,y_test, estimator,name="Estimator"):
    print("-----Metrics of {} estimator-----\n".
    ↪format(name))
    print(classification_report(y_test, estimator.predict(X_test)))
    plot_confusion_matrix(X_test,y_test, estimator,name)
    plot_roc(X_test,y_test,estimator,name)
```

```
[24]: display_metrics(X_test,y_test, logreg.best_estimator_, "Logistic Regression_
    ↪with L2 Norm")
```

```
-----Metrics of Logistic Regression with L2 Norm
estimator-----
```

	precision	recall	f1-score	support
0	1.00	0.82	0.90	28
1	0.87	1.00	0.93	33
accuracy			0.92	61
macro avg	0.93	0.91	0.92	61
weighted avg	0.93	0.92	0.92	61




```
[25]: param_grid = {
        'C':np.append(np.linspace(0.1,1, num=10),[1]),
        'kernel':['linear', 'poly', 'rbf', 'sigmoid'],
        'gamma':['scale', 'auto'],
        'probability':[True],
        'random_state': [123]
    }
    # instantiate the grid
    grid = GridSearchCV(SVC(max_iter=1000), param_grid, cv=10, scoring='accuracy',
        ↪return_train_score=True)
    SVC_est = grid.fit(X_train,y_train)
```

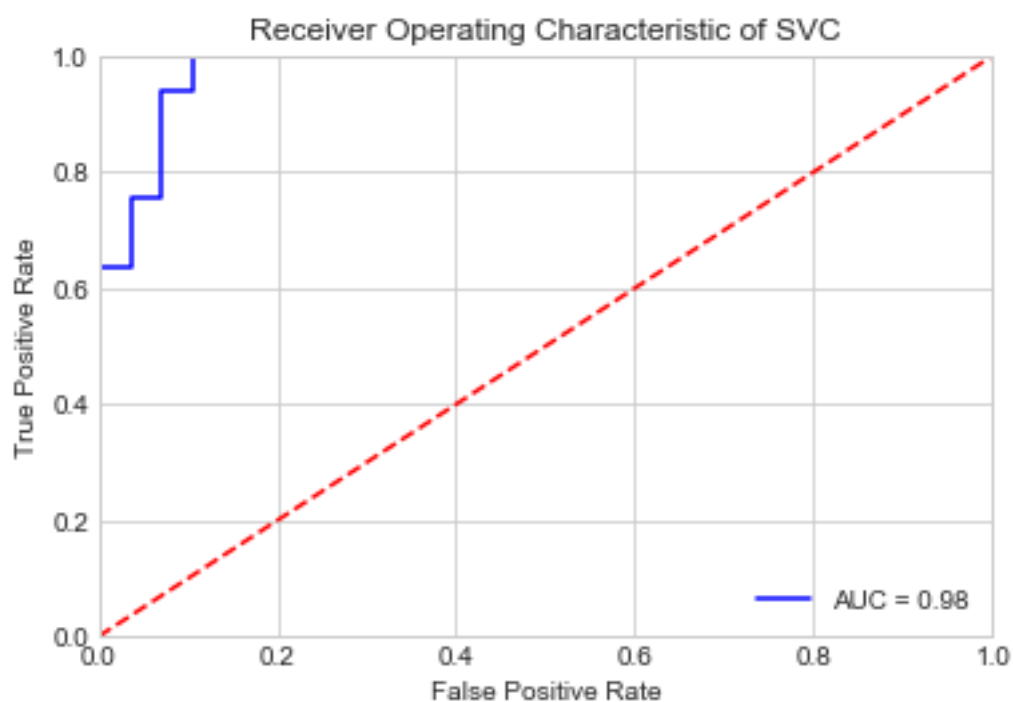
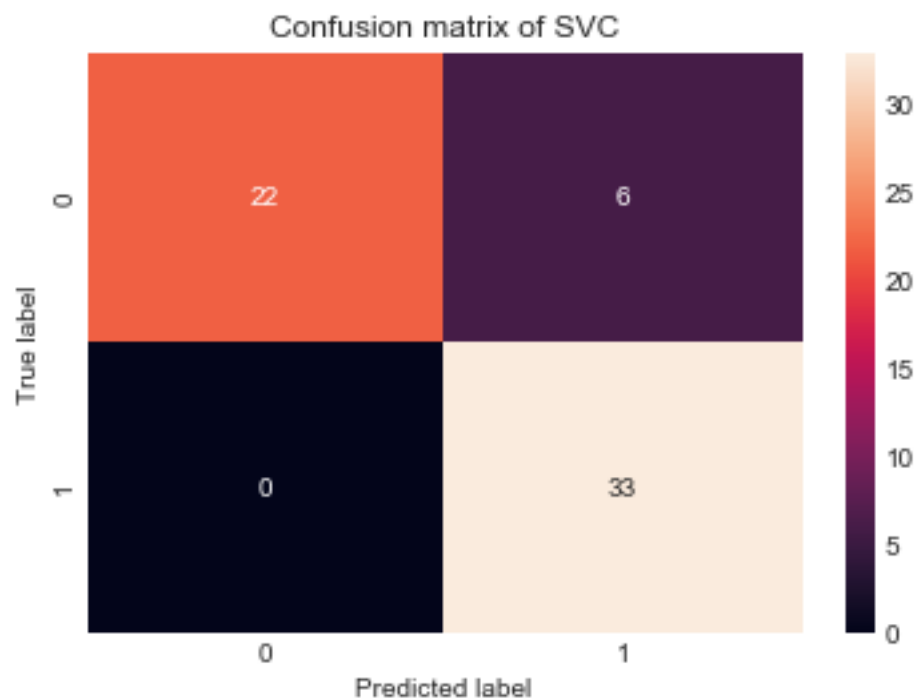
```
[26]: SVC_est.best_estimator_
```

```
[26]: SVC(C=0.8, break_ties=False, cache_size=200, class_weight=None, coef0=0.0,
        decision_function_shape='ovr', degree=3, gamma='scale', kernel='sigmoid',
        max_iter=1000, probability=True, random_state=123, shrinking=True,
        tol=0.001, verbose=False)
```

```
[27]: display_metrics(X_test,y_test, SVC_est.best_estimator_, "SVC")
```

```
-----Metrics of SVC estimator-----
```

	precision	recall	f1-score	support
0	1.00	0.79	0.88	28
1	0.85	1.00	0.92	33
accuracy			0.90	61
macro avg	0.92	0.89	0.90	61
weighted avg	0.92	0.90	0.90	61




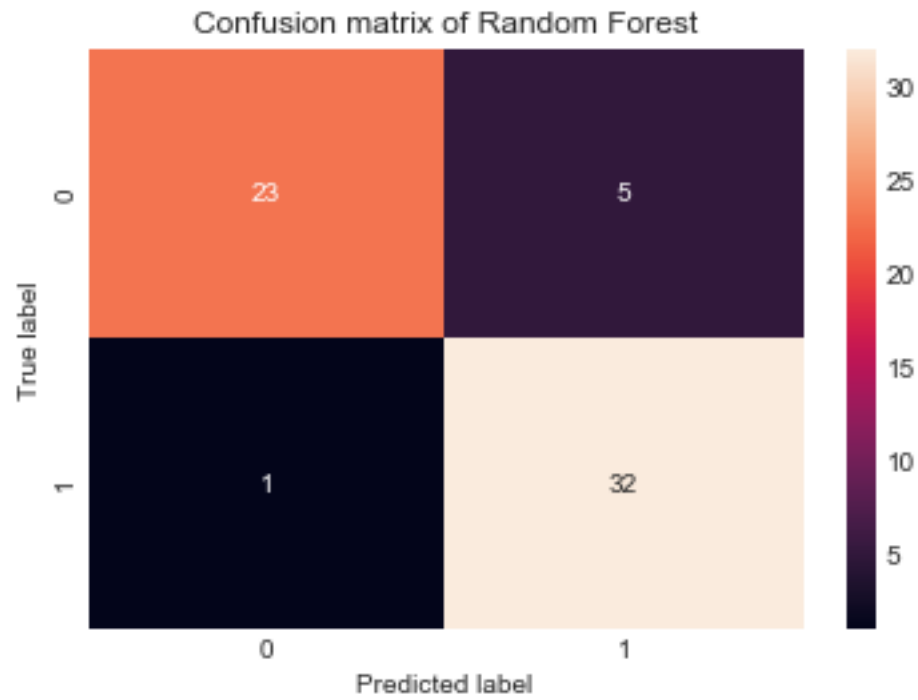
```
[28]: param_grid = {
        'n_estimators': np.linspace(50,150, num=50).astype(int),
        'criterion': ['gini', 'entropy'],
        'max_features': ['auto', 'log2'],
        'random_state': [123]
    }
    # instantiate the grid
    grid = GridSearchCV(RandomForestClassifier(), param_grid, cv=10,
        ↳scoring='accuracy', return_train_score=True)
    rf_est = grid.fit(X_train,y_train)
    rf_est.best_estimator_
```

```
[28]: RandomForestClassifier(bootstrap=True, ccp_alpha=0.0, class_weight=None,
                             criterion='entropy', max_depth=None, max_features='auto',
                             max_leaf_nodes=None, max_samples=None,
                             min_impurity_decrease=0.0, min_impurity_split=None,
                             min_samples_leaf=1, min_samples_split=2,
                             min_weight_fraction_leaf=0.0, n_estimators=50,
                             n_jobs=None, oob_score=False, random_state=123,
                             verbose=0, warm_start=False)
```

```
[29]: display_metrics(X_test,y_test,rf_est.best_estimator_,"Random Forest")
```

-----Metrics of Random Forest estimator-----

	precision	recall	f1-score	support
0	0.96	0.82	0.88	28
1	0.86	0.97	0.91	33
accuracy			0.90	61
macro avg	0.91	0.90	0.90	61
weighted avg	0.91	0.90	0.90	61



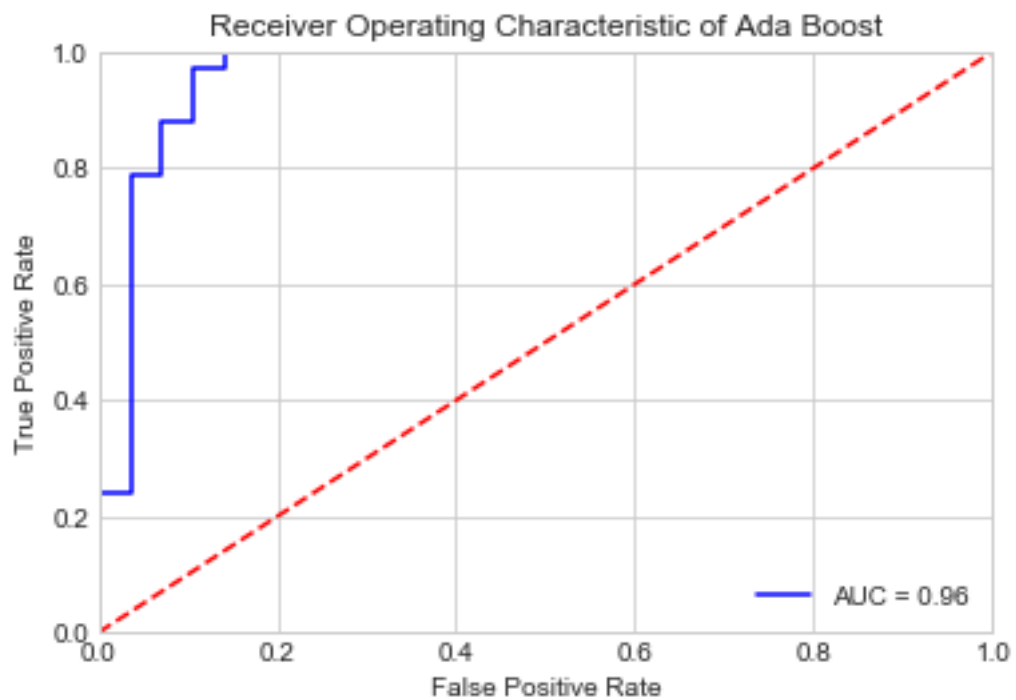
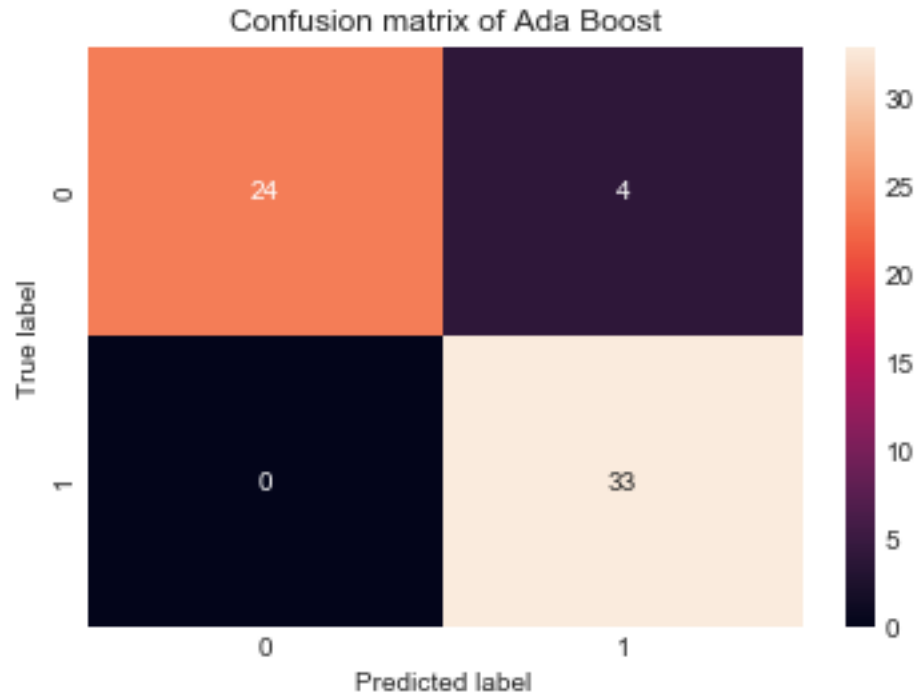
```
[35]: param_grid = {
        'base_estimator': [None, SVC(max_iter=1000),
        ↳ LogisticRegression(max_iter=1000), logreg.best_estimator_],
        'n_estimators': np.linspace(50,150, num=50).astype(int),
        'random_state': [123]
    }
    # instantiate the grid
    grid = GridSearchCV(AdaBoostClassifier(), param_grid, cv=10,
    ↳ scoring='accuracy', return_train_score=True)
    ada_est = grid.fit(X_train,y_train)
    ada_est.best_estimator_
```

```
[35]: AdaBoostClassifier(algorithm='SAMME.R',
                        base_estimator=LogisticRegression(C=1.0, class_weight=None,
                                                            dual=False,
                                                            fit_intercept=True,
                                                            intercept_scaling=1,
                                                            l1_ratio=None,
                                                            max_iter=1000,
                                                            multi_class='auto',
                                                            n_jobs=None, penalty='l2',
                                                            random_state=None,
                                                            solver='lbfgs', tol=0.0001,
                                                            verbose=0,
                                                            warm_start=False),
                        learning_rate=1.0, n_estimators=147, random_state=123)
```

```
[34]: display_metrics(X_test,y_test,ada_est.best_estimator_,"Ada Boost")
```

-----Metrics of Ada Boost estimator-----

	precision	recall	f1-score	support
0	1.00	0.86	0.92	28
1	0.89	1.00	0.94	33
accuracy			0.93	61
macro avg	0.95	0.93	0.93	61
weighted avg	0.94	0.93	0.93	61



From each of the classification reports we can see that AdaBoost performed the best with an

accuracy of 93% and a true positive accuracy of 100%. In this use case having false positives aren't as bad. Having true negatives is very bad because it means there is a misdiagnosis which can lead to death. In this case 93% accuracy is not good enough for proper medical diagnosis but with a low true negative error this classifier can be used to do a preliminary check on patients. There are many ways the model can be improved. The use of more data can vastly improve the model, XGBoost and possibly deep learning can easily improve the performance of the model. The use of more powerful models will be worth the computational cost because it will decrease the classification error.

0.5 Save the model

```
[33]: pickle.dump(ada_est.best_estimator_, open('model.pkl', 'wb'))
```

0.6 Reference

The dataset was provided by the UCI Machine Learning Repository with the help of: 1. Hungarian Institute of Cardiology. Budapest: Andras Janosi, M.D. 2. University Hospital, Zurich, Switzerland: William Steinbrunn, M.D. 3. University Hospital, Basel, Switzerland: Matthias Pfisterer, M.D. 4. V.A. Medical Center, Long Beach and Cleveland Clinic Foundation: Robert Detrano, M.D., Ph.D.